

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
In [2]: link = 'https://raw.githubusercontent.com/AdiPersonalWorks/Random/master/student_scores%20-%20student_scores.csv'
df = pd.read_csv(link)
print("successful import")
```

successful import

```
In [3]: df.head(5)
```

```
Out[3]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [4]: df.head()
```

```
Out[4]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [5]: df.tail()
```

```
Out[5]:
```

	Hours	Scores
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

```
In [6]: df.columns
```

```
Out[6]: Index(['Hours', 'Scores'], dtype='object')
```

```
In [7]: df.shape
```

```
Out[7]: (25, 2)
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 25 entries, 0 to 24  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  ---  
0   Hours    25 non-null      float64  
1   Scores   25 non-null      int64  
dtypes: float64(1), int64(1)  
memory usage: 528.0 bytes
```

```
In [9]: df.describe()
```

```
Out[9]:
```


	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
In [10]: df.corr()
```

```
Out[10]:
```


	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

```
In [11]: # Plotting the distribution of scores  
df.plot(x='Hours', y='Scores', style='o')  
plt.title('Hours vs Percentage')  
plt.xlabel('Hours Studied')  
plt.ylabel('Percentage Score')  
plt.show()
```



```
In [12]: sns.jointplot(y='Hours',x='Scores', data= df)
```

```
Out[12]: <seaborn.axisgrid.JointGrid at 0x16a77b6a100>
```



### Data Preprocessing

```
In [13]: x = df.iloc[:, :-1].values  
y = df.iloc[:, 1].values
```

### Model Training

```
In [14]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)  
model = LinearRegression()  
model.fit(x_train.reshape(-1,1), y_train)  
print('done')  
  
done
```

### Predictions

```
In [15]: # Testing data  
print(x_test)  
# Model Prediction  
y_pred = model.predict(x_test)  
  
[[1.5]  
 [3.2]  
 [7.4]  
 [2.5]  
 [5.9]]
```

```
In [16]: newDF = pd.DataFrame({'Reg':y_test , 'Predicted':y_pred})  
newDF
```

```
Out[16]:
```

	Reg	Predicted
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

```
In [17]: #Estimating training and test score
print("Training Score:",model.score(x_train,y_train)*100,'%')
```


Training Score: 95.15510725211553 %

```
In [18]: print("Test Score:",model.score(x_test,y_test) * 100,'%')
```

Test Score: 94.54906892105356 %

```
In [19]: # Plotting the Bar graph to depict the difference between the actual and predicted value
```

```
newDF.plot(kind='bar',figsize=(5,5))
plt.grid(which='major', linewidth='0.5', color='black')
plt.grid(which='minor', linewidth='0.5', color='black')
plt.show()
```



```
In [20]: # testing the given sample
```

```
hours = 9.25
newTest=np.array([hours])
newTest = newTest.reshape(-1,1)
newPred = model.predict(newTest)
```

```
In [21]: print("No of Hours = {}".format(hours))
print("Predicted Score = {}".format(newPred[0]))
```

No of Hours = 9.25  
Predicted Score = 93.69173248737538

```
In [22]: print('Mean Absolute Error:',metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R-2:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 4.183859899002975  
Mean Squared Error: 21.5987693072174  
Root Mean Squared Error: 4.6474476121003665  
R-2: 0.9454906892105356